

# Stat 432 Homework 4

Assigned: Sep 20, 2018; Due: Sep 28, 2018


## Question 1 (more on PCA)

[3 points] Take the first four columns of the `iris` data to verify the connection of PCS with the singular value decomposition (`svd`).

Load `iris` data. We use `svd` function to obtain the singular values  $\mathbf{D}$ . Then PCA variance is the eigenvalues of  $\hat{\Sigma} = \mathbf{X}^T \mathbf{X} / (n - 1)$ , which is  $D^2 / (n - 1)$ , and we plot them in the decreasing order. (1 point)


```
# load 'iris'
data(iris)
n=nrow(iris)
# SVD
irissvd=svd(scale(iris[, 1:4], center = TRUE, scale = FALSE))
plot(irissvd$d^2/(n-1), type = 'b', pch = 19, xlab='', ylab='Variances', main = "Iris PCA Variance")
```

**Iris PCA Variance**



Now we obtain the principal components by  $\mathbf{X}^T \mathbf{V}$  and plot the first two on a figure with colors associated with species. (1 point) And we print the rotation matrix  $\mathbf{V}$ . (1 point)

```
# principal components
irispc=as.matrix(iris[, 1:4])%*%irissvd$v
plot(irispc[,1], irispc[,2], pch=20, col=c(1:3)[iris$Species], xlab=expression(pc[1]), ylab=expression(pc[2]), legend='topright', pch=20, legend=unique(iris$Species), col=c(1:3))
```



```
# rotation matrix
irissvd$v
```

```
##           [,1]      [,2]      [,3]      [,4]
## [1,]  0.36138659 -0.65658877  0.58202985  0.3154872
## [2,] -0.08452251 -0.73016143 -0.59791083 -0.3197231
## [3,]  0.85667061  0.17337266 -0.07623608 -0.4798390
## [4,]  0.35828920  0.07548102 -0.54583143  0.7536574
```

Question 2 ( $k$ -NN for classification)

[4 points] Consider again the zip code digits data. And we will use the Euclidean distance. We want to predict the digit of the 4th observation in the testing dataset.

```
library(ElemStatLearn)
train.x = zip.train[, -1]
train.y = as.factor(zip.train[, 1])
test.x.one = zip.test[4, -1]
```

One way to find 15 nearest neighbors using `sweep` and `rowSums` could be as follows: (1 point)

```
k=15
# find k nn
(idx_knn=which(rank(rowSums((sweep(train.x, 2, test.x.one, FUN = "-")^2)) <= k)))
```

```
## [1] 389 521 1619 1825 2240 3471 3988 4187 4188 5106 5143 5198 6450 6774
## [15] 6976
```

Then we can find the most frequent digit among these 15 observations. We find out that the majority vote is "0", not the true label "6". (1 point)

```
# majority vote
(major_vote=names(which.max(table(train.y[idx_knn]))))
```

```
## [1] "0"
```

```
# true label
(test.y.one=zip.test[4,1])
```

```
## [1] 6
```

Now we repeat the process for a sequence of  $k$  and list the results as follows (1 point. No need to list for multiple  $k$ 's for this step; you get full credit as long as you get correct prediction for some  $k$ .)

```
# try different k's
pred=rep(NA,15)
for(k in 1:15){
  idx_knn=which(rank(rowSums((sweep(train.x, 2, test.x.one, FUN = "-"))^2)) <= k)
  major_vote=names(which.max(table(train.y[idx_knn])))
  pred[k]=as.numeric(major_vote)==test.y.one
}
# print out results
names(pred)<-1:15
pred
```

```
##      1      2      3      4      5      6      7      8      9     10     11     12
## TRUE  TRUE  TRUE FALSE  TRUE  TRUE  TRUE  TRUE  TRUE  TRUE  TRUE FALSE
##     13     14     15
## FALSE FALSE FALSE
```

For this given testing case, we see that with small  $k$ , it happens to predict the correct labels though it fails at  $k = 4$ . With increasing  $k$ , it gets correct prediction until  $k > 10$ . Now we repeat the process for the first 100 observations with  $k$  ranging from 1 to 20, and plot the error rate wrt  $k$ .

```
# pairwise distance function
# cited from https://www.r-bloggers.com/pairwise-distances-in-r/
pdist <- function(A,B) {
  an = apply(A, 1, function(rvec) crossprod(rvec,rvec))
  bn = apply(B, 1, function(rvec) crossprod(rvec,rvec))


  m = nrow(A)
  n = nrow(B)

  tmp = matrix(rep(an, n), nrow=m)
  tmp = tmp + matrix(rep(bn, m), nrow=m, byrow=TRUE)
  # sqrt( tmp - 2 * tcrossprod(A,B) )
  tmp - 2 * tcrossprod(A,B)
}

# obtain the first 100 observations as testing data
test.x.hundred = zip.test[1:100, -1]
test.y.hundred = zip.test[1:100, 1]
# repeat for k from 1 to 20
errate=rep(NA,20)
dist_trte=pdist(train.x,test.x.hundred)
for(k in 1:20){
  idx_knn=apply(dist_trte,2,function(d) which(rank(d)<=k))
  if(k==1){idx_knn=t(idx_knn)}
  major_vote=apply(idx_knn,2,function(id)names(which.max(table(train.y[id]))))
  errate[k]=mean(as.numeric(major_vote)!=test.y.hundred)
}

# plot the error rate
plot(errate,type = 'b', pch = 19, xlab =expression(k), ylab= 'Prediction Error Rate', main = "Bias-Variation")
```

## Bias–Variance Trade–Off



Seen from the results, the best prediction result (lowest error rate) is attained at  $k = 3$ . (1 point. You can also come up other criterion for selecting the best  $k$ . You lose 0.5 points if not reporting the best  $k$  as the final answer.)

Question 3 (Cross-validation using the `caret` package)

[3 points]

Install and load the `caret` package. We use it to do CV (1 point) to choose the best  $k$ . And then we test the best fit on testing data.

```
# load 'caret' which requires the package 'e1071' that needs to be installed
library(caret)
# prepare for the training
TrainData = data.frame(x = train.x)
TrainClasses = as.factor(train.y)
# train the knn model for classification using CV, taking several minutes
knnFit <- train(TrainData, TrainClasses, method = "knn", tuneGrid = data.frame("k" = c(1:10)),
               trControl = trainControl(method = "cv", number = 3))
# best k
knnFit$bestTune

##      k
## 1 1

# prepare for the testing
TestData = data.frame(x = zip.test[, -1])
TestClasses = as.factor(zip.test[, 1])
# test
# predy=extractPrediction(list(knn=knnFit), testX = TestData, testY = TestClasses, unkX = TRUE)
predy=predict(knnFit, newdata=TestData)
# confusion matrix
table(predy, TestClasses)

##      TestClasses
## predy  0  1  2  3  4  5  6  7  8  9
##      0 355  0  6  3  0  2  0  0  5  0
```

##	1	0	255	1	0	3	1	0	1	0	0
##	2	2	0	183	2	1	2	1	1	1	1
##	3	0	0	2	154	0	4	0	1	6	0
##	4	0	6	1	0	182	0	2	4	1	2
##	5	0	0	0	5	1	145	3	0	1	0
##	6	0	2	0	0	2	2	164	0	0	0
##	7	1	1	2	0	2	0	0	139	1	4
##	8	0	0	3	0	1	3	0	0	148	1
##	9	1	0	0	2	8	1	0	1	3	169

The best  $k = 1$  (1 point), which usually tends to overfit data. This may be due to the small number folds. The confusion matrix (1 point) shows it works well in prediction.