

STAT 432: Basics of Statistical Learning

Introduction to Convex Optimization

Shiwei Lan, Ph.D. <shiwei@illinois.edu>

<http://shiwei.stat.illinois.edu/stat432.html>

January 18, 2019

University of Illinois at Urbana-Champaign

Convex Optimization

- This lecture gives a very brief introduction to convex optimization
- The goal is to have sufficient knowledge to deal with specific problems such as Lasso, SVM, etc.
- Reference:
Boyd, Stephen, and Lieven Vandenberghe. *Convex Optimization*. Cambridge University Press, 2004.
- Many of the figures in this lecture are taken from online sources. I want to thank all of them!

- The problem: minimizing a convex function in a convex set

$$\begin{aligned} & \underset{\beta}{\text{minimize}} && f(\beta) \\ & \text{subject to} && g_i(\beta) \leq 0, \quad i = 1, \dots, m \\ & && \mathbf{A}\beta = b \end{aligned}$$



- Examples:
 - Linear regression: minimize $\frac{1}{2} \|\mathbf{y} - \mathbf{X}\beta\|^2$, subject to none.
 - Ridge regression: minimize $\frac{1}{2} \|\mathbf{y} - \mathbf{X}\beta\|^2$, subject to $\sum_{j=1}^p \beta_j^2 < s$
 - First principal component: maximize $\beta^\top \mathbf{X}^\top \mathbf{X} \beta$, subject to $\beta^\top \beta = 1$

Convex Optimization

- What is a **convex set** $C \in \mathbb{R}^p$?

$$\beta_1, \beta_2 \in C \implies \alpha\beta_1 + (1 - \alpha)\beta_2 \in C, \quad \forall 0 \leq \alpha \leq 1.$$

- Visual:




Convex Optimization

- What is a **convex function** $f : \mathbb{R}^p \rightarrow \mathbb{R}$?

$$f(\alpha\beta_1 + (1 - \alpha)\beta_2) \leq \alpha f(\beta_1) + (1 - \alpha)f(\beta_2) \quad \forall 0 \leq \alpha \leq 1.$$

- In Probability: Jensen's inequality
- Visual:



Convex functions

- To comply with notations in the literature, I will use \mathbf{x} as the argument instead of using β , and we are interested in the function $f(\mathbf{x})$.
- Examples of convex functions:
 - $\exp(x)$, $-\log(x)$, etc.
 - Affine: $a^\top \mathbf{x} + b$ is both convex and concave
 - Quadratic: $\frac{1}{2} \mathbf{x}^\top \mathbf{A} \mathbf{x} + b^\top \mathbf{x} + c$, if \mathbf{A} is positive semidefinite.
 - All norms: ℓ_p
- A function is strictly convex if we can remove the equal sign:

$$f(\alpha \beta_1 + (1 - \alpha) \beta_2) < \alpha f(\beta_1) + (1 - \alpha) f(\beta_2) \quad \forall 0 \leq \alpha \leq 1.$$

- f is convex $\iff -f$ is concave

Properties of Convex functions

- **First-order property:** If f is differentiable with convex domain, then f is convex iff

$$f(\mathbf{x}^*) \geq f(\mathbf{x}) + \nabla f(\mathbf{x})^\top (\mathbf{x}^* - \mathbf{x})$$

- If we have a feasible point \mathbf{x} with $\nabla f(\mathbf{x}) = \mathbf{0}$, it means all alternative points \mathbf{x}^* have larger function value $f(\mathbf{x}^*) \geq f(\mathbf{x})$.
- Hence, we call \mathbf{x} a local minimizer. **It may not be unique**, but its as good as any other solution.
- Example: In a linear regression if we have linearly dependent columns in the design matrix. The solution of parameters is not unique.

Properties of Convex functions

- **Second-order property:** If f is twice differentiable with convex domain, then f is convex iff

$$\nabla^2 f(\mathbf{x}) \succeq 0 \quad \text{for any } \mathbf{x} \text{ in the domain,}$$

where

$$\mathbf{H}(\mathbf{x}) = \nabla^2 f(\mathbf{x}) = \left(\frac{\partial^2 f(\mathbf{x})}{\partial x_i \partial x_j} \right)$$

- $\mathbf{H}(\mathbf{x})$ is the Hessian matrix.
- If $\nabla^2 f(\mathbf{x}) \succ 0$ (positive definite), meaning f is strictly convex, then a local minimizer is also a global minimizer.
- Example: In linear regression when $\mathbf{X}^\top \mathbf{X} \succ 0$, i.e., invertible.

Solving convex problems

- In many situations, we just deal with an unconstrained, smooth convex function

$$\underset{\beta}{\text{minimize}} \quad f(\beta)$$

- OLS and ridge, etc. can all be formulated as this problem.
- Lasso is not smooth at some particular points of β

Gradient descent

- One of the simplest algorithm is gradient descent.
- At any given point \mathbf{x} , we want to move it to the direction where f can decrease.
- Consider the Taylor expansion near \mathbf{x} :

$$\begin{aligned} f(\mathbf{x}^*) \\ \approx f(\mathbf{x}) + \nabla f(\mathbf{x})^\top (\mathbf{x}^* - \mathbf{x}) + \frac{1}{2} (\mathbf{x}^* - \mathbf{x})^\top \mathbf{H}(\mathbf{x}) (\mathbf{x}^* - \mathbf{x}) \end{aligned}$$

- If we minimize this quadratic approximation, the new point that gives the smallest $f(\mathbf{x}^*)$ is

$$\mathbf{x}^* = \mathbf{x} - \mathbf{H}(\mathbf{x})^{-1} \nabla f(\mathbf{x})$$



Gradient descent

- Since calculating the second derivative (and inverse) $\mathbf{H}(\mathbf{x})$ can be difficult, let's just use an identity matrix $\frac{1}{\delta} \mathbf{I}$.
- Then the new point is

$$\mathbf{x}^* = \mathbf{x} - \delta \nabla f(\mathbf{x})$$

- Gradient descent uses this updating scheme to iteratively archive smaller $f(\mathbf{x})$.
- However, we have to choose δ , which is known as the step size.
 - A step size **too large** may not even converge at all.
 - How about we just fix δ to be a small value, say 10^{-5} .
 - A step size **too small** will take many iterations to converge.
 - Line search is usually used. Sometimes, inexact line search.

Gradient descent



Newton-Raphson

- When we **have explicit formula of the Hessian**, we can simply compute them at each point \mathbf{x} and follow the updating scheme

$$\mathbf{x}^{(k+1)} = \mathbf{x}^{(k)} - \mathbf{H}(\mathbf{x}^{(k)})^{-1} \nabla f(\mathbf{x}^{(k)})$$

or, sometimes for numerical stability,

$$\mathbf{x}^{(k+1)} = \mathbf{x}^{(k)} - \delta \mathbf{H}(\mathbf{x}^{(k)})^{-1} \nabla f(\mathbf{x}^{(k)})$$

- However, this is rare in practice since most objective functions are very complicated. Constraints may also create difficulties.
- We can do this for ridge, but not Lasso.

Numerical approximation

- Sometimes even computing the gradient $\nabla f(\mathbf{x})$ is difficult.
- We can simply approximate them numerically (element-wise) by the definition

$$\nabla f(\mathbf{x})_j \approx \frac{f(\mathbf{x} + \delta \mathbf{e}_j) - f(\mathbf{x})}{\delta}$$

where \mathbf{e}_j is a vector with the j th element 1 and 0 everywhere else.

- Numerically approximating the gradient could be slow, depending on the size of the problem.

Numerical approximation

- In many cases, we have exact formula of the gradient, but want to speed things up with the second order information.
- You usually don't want to directly numerically approximate $\mathbf{H}(\mathbf{x})$ because its very very expensive.
- However, quasi-Newton methods, such as BFGS, progressively approximates $\mathbf{H}(\mathbf{x})^{-1}$ by utilizing the Sherman-Morrison formula.

- Gradient descent can be viewed as using $\mathbf{H}(\mathbf{x})^{-1} = \mathbf{I}$
- $\mathbf{x}^{(0)} \rightarrow \mathbf{x}^{(1)} \rightarrow \mathbf{x}^{(2)} \rightarrow \dots \rightarrow \mathbf{x}^{(k)} \rightarrow \mathbf{x}^{(k+1)} \rightarrow$
- Along this path, we computed $\rightarrow \nabla f(\mathbf{x}^{(k)}) \rightarrow \nabla f(\mathbf{x}^{(k+1)}) \rightarrow \dots$
- If we treat the function $f(\mathbf{x}^{(k)})$ locally as a quadratic function,

$$\nabla f(\mathbf{x}^{(k+1)}) - \nabla f(\mathbf{x}^{(k)}) \approx \mathbf{H}(\mathbf{x}^{(k)})(\mathbf{x}^{(k+1)} - \mathbf{x}^{(k)})$$

- This leads to rank one/two updates of \mathbf{H} . Read [DFP](#) and [BFGS](#) references.

Broyden, Fletcher, Goldfarb, Shanno



Implementation

- If you have a smooth objective function, usually the log-likelihood $L(\mathbf{y}, \mathbf{X}, \beta)$, and we want to solve the parameters β .
- You can utilize the `optim()` function in R

```
1 > L <- function(b, X, Y) ...  
2 > bhat = optim(rep(0, P), L, X = X, Y = Y, method = "BFGS")
```

- Sometimes, using a different initial value may be better.

None differentiable problems

- Non-differentiable problems become prevalent when we add penalties to the objective function
- Usually, these problems (at least the ones that we care about) are decomposable, meaning that

$$f(\mathbf{x}) = g(\mathbf{x}) + h(\mathbf{x})$$

with differentiable g and non-differentiable h but still convex. For example:

$$\frac{1}{2} \|\mathbf{y} - \mathbf{X}\boldsymbol{\beta}\|_2^2 + \lambda \|\boldsymbol{\beta}\|_1$$

- There are other general approaches that can deal with this problem, but we will only look at the special case: Lasso.

- The Lasso problem has a special form, i.e., its **separable**:

$$f(\mathbf{x}) = g(\mathbf{x}) + \sum_{i=1}^p h(x_i)$$

- I will switch back to the “ β ” notation for parameters:

$$f(\beta) = \frac{1}{2}(\mathbf{y} - \mathbf{X}\beta)^\top(\mathbf{y} - \mathbf{X}\beta) + \sum_{i=1}^p |\beta_i|$$

- Instead of updating all parameters each step, we only update one parameter each step and loop over all parameters.

Coordinate descent

- The **Gauss-Seidel style** coordinate descent algorithm goes like, at the k th (grand) iteration:

$$\beta_1^{(k+1)} = \arg \min_{\beta_1} f(\beta_1, \beta_2^{(k)}, \dots, \beta_p^{(k)})$$

$$\beta_2^{(k+1)} = \arg \min_{\beta_2} f(\beta_1^{(k+1)}, \beta_2, \dots, \beta_p^{(k)})$$

...

$$\beta_p^{(k+1)} = \arg \min_{\beta_p} f(\beta_1^{(k+1)}, \beta_2^{(k+1)}, \dots, \beta_p)$$

- After we complete this loop, all β_j are updated to their new values, and we start over.

Coordinate descent

- The **Jacobi style** algorithm goes like, at the k th (grand) iteration:

$$\beta_1^{(k+1)} = \arg \min_{\beta_1} f(\beta_1, \beta_2^{(k)}, \dots, \beta_p^{(k)})$$

$$\beta_2^{(k+1)} = \arg \min_{\beta_2} f(\beta_1^{(k)}, \beta_2, \dots, \beta_p^{(k)})$$

...

$$\beta_p^{(k+1)} = \arg \min_{\beta_p} f(\beta_1^{(k)}, \beta_2^{(k)}, \dots, \beta_p)$$

- After we complete this loop, update all β_j to their new values, and start over.
- Jacobi style algorithm can be computed in a **parallel** fashion, while Gauss-Seidel style can only be done **sequentially**.

Coordinate descent

- Two questions:
 - 1) is this going to be slower or faster than gradient descent?
 - 2) will it converge?
- Lets take a linear regression as an example (no penalty):

$$f(\beta) = \frac{1}{2}(\mathbf{y} - \mathbf{X}\beta)^\top(\mathbf{y} - \mathbf{X}\beta)$$

- At each iteration, for each β_i , lets fix all other parameters $\beta_{(-j)}$.
- Suppose we do not know the Hessian matrix, the gradient descent goes like

$$\begin{aligned}\beta &= \beta - \delta \nabla f(\beta) \\ &= \beta - \delta \mathbf{X}^\top(\mathbf{y} - \mathbf{X}\beta)\end{aligned}$$

- This cost $O(np)$ flops.

Coordinate descent

- What about coordinate descent? It is a one-variable regression problem if we fix $\beta_{(-j)}$:

$$f(\beta_j^{(k+1)}) = \frac{1}{2} \|\mathbf{y} - X_j \beta_j - \mathbf{X}_{(-j)} \beta_{(-j)}^{(k)}\|_2^2$$

- Define $\mathbf{r} = \mathbf{y} - \mathbf{X}_{(-j)} \beta_{(-j)}^{(k)} = \mathbf{y} - \mathbf{X} \beta^{(k)} + X_j \beta_j^{(k)}$


$$\beta_j^{(k+1)} = \frac{\mathbf{X}_j^\top \mathbf{r}}{\mathbf{X}_j^\top \mathbf{X}_j} = \frac{\mathbf{X}_j^\top (\mathbf{y} - \mathbf{X} \beta^{(k)})}{\mathbf{X}_j^\top \mathbf{X}_j} + \beta_j^{(k)}$$

- After updating β_j , put $X_j \beta_j^{(k+1)}$ back into \mathbf{r} , then subtract the effect from $(j+1)$ for the next update
- Usually this is proceeded with the Gauss-Seidel style.

Coordinate descent

- **Intuition:** we firstly take out all we currently explained ($\mathbf{X}\beta$) from \mathbf{y} , save that as the residual, \mathbf{r} . Then we look in turn, how much extra each variable can explain.
- Updating each β_j cost $O(n)$ flops, then each iteration cost $O(np)$, same as gradient descent.
- What about their performances?


Coordinate descent



Coordinate descent

- What happened?
- Its not really a fair game: gradient descent only utilize the first order information, while coordinate descent updated each coordinate fully, at least for each iteration.
- When is coordinate descent useful/better?
 - If updating each coordinate is cheap, and maybe the solution is explicit (our lasso problem will be of this type).
 - More importantly, the problem has to be separable
- When will coordinate descent fail?

Coordinate descent



Coordinate descent for Lasso

- The Lasso solution can be obtained in the same way.
- Recall that the penalty is coordinate-wise, for each update, we are solving

$$f(\beta_j^{(k+1)}) = \frac{1}{2} \|\mathbf{y} - X_j \beta_j - \mathbf{X}_{(-j)} \boldsymbol{\beta}_{(-j)}^{(k)}\|_2^2 + \lambda |\beta_j|$$

- Again, this is a one-variable optimization problem, and we have the exact solution (see lecture note “Penalized”), which is a soft-threshold version of the OLS estimator.
- A **path-wise algorithm** starts with a large λ value, and run until no β can be moved anymore (converged), then reduce the λ by a factor and start over the update.
- This is how the **glmnet** package solves Lasso.