

# STAT 432: Basics of Statistical Learning

## Clustering and PCA

---

Shiwei Lan, Ph.D. <[shiwei@illinois.edu](mailto:shiwei@illinois.edu)>

<http://shiwei.stat.illinois.edu/stat432.html>

January 22, 2019

University of Illinois at Urbana-Champaign

# Unsupervised Learning

---

# Unsupervised Learning

- No response variable  $Y$ , only  $\{x_i\}_{i=1}^n$ .
- Goal: learn patterns in  $X$ .
- Examples
  - Estimate the **density**, **covariance**, **graph (network)**, etc. of  $X$  — could be difficult in high-dimensional settings
  - **Cluster analysis**: identify multiple regions of the feature space that contains modes of density.
  - **Dimension reduction**: identify low-dimensional manifold within the feature space  $\mathcal{X}$  that represents high data density.
- Often times, there is no clear measure of success.

# Cluster Analysis

---

# Cluster Analysis

- Group the dataset into subsets so that those within each subset are more closely related (similar) to each other than those objects assigned to other subsets. Each subset is called a cluster
- Flat clustering vs. hierarchical clustering: flat clustering divides the dataset into  $k$  cluster, and hierarchical clustering arranges the clusters into a natural hierarchy.
- Clustering results are crucially dependent on the measure of similarity (or distance) between the “points” to be clustered.

# Distance Metric

- A **distance metric** or a distance function is a function that defines the similarity of two elements (points, sets, etc.)
- For the distance of two points (with continuous entries), the most commonly used measurement is the **Euclidian distance**:

$$\begin{aligned}d(u, v) &= \|u - v\|_2 \\&= \sqrt{\sum_{j=1}^p (u_j - v_j)^2}\end{aligned}$$

- For categorical entries, the **Hamming distance** is usually used

$$d(u, v) = \sum_{j=1}^p \mathbf{1}\{u_j \neq v_j\}$$

- Distance measures should be defined based on the application. There is no universally best approach.

# Clustering

- Suppose we have a set of  $n$  data points
- We want to form  $K \ll n$  clusters, indexed by  $k \in \{1, \dots, K\}$ .
- Let  $C(\cdot)$  be a cluster index function that assign th  $i$ th observation or cluster  $C(i)$ .
- Consider: search for a function  $C : \{1, \dots, n\} \rightarrow \{1, \dots, K\}$  to minimize the within cluster distance:

$$W(C) = \frac{1}{2} \sum_{k=1}^K \sum_{C(i), C(i')=k} d(x_i, x_{i'}).$$

# Clustering

- This is equivalent to maximizing the **between cluster distance**

$$B(C) = \frac{1}{2} \sum_{k=1}^K \sum_{C(i)=k} \sum_{C(i') \neq k} d_{ii'}$$

- Note that the **total distance** can be broke down into

$$\begin{aligned} T &= \frac{1}{2} \sum_{i=1}^n \sum_{i'=1}^n d_{ii'} = \frac{1}{2} \sum_{k=1}^K \sum_{C(i)=k} \left[ \sum_{C(i')=k} d_{ii'} + \sum_{C(i') \neq k} d_{ii'} \right] \\ &= W(C) + B(C) \end{aligned}$$

- The total distance is fixed for a given set of data, hence

$$\text{minimize } W(C) \iff \text{maximize } B(C)$$



- Given a specific distance measure  $d(\cdot, \cdot)$ , several algorithms can be used to find the clusters
  - Combinatorial algorithm
  - $K$ -means clustering
  - Hierarchical clustering

# Combinatorial Algorithm

---

# Combinatorial Algorithms

- For **small**  $n$  and  $K$ , we could minimize  $W$  by **brute-force** search.
- However, the number of “tries” needed to complete the search is

$$S(n, K) = \frac{1}{K!} \sum_{k=1}^K (-1)^{K-k} \binom{K}{k} k^n$$

- For example  $S(10, 4) = 34,105$ ;  $S(19, 4) \approx 10^{10}$ .
- This is **not feasible for large  $n$  and  $K$** , since the number of distinct assignments can be extremely large.
- It calls for more efficient algorithms: **may not be optimal** but a reasonably good suboptimal partition.

# *K*-means Clustering

---

# K-means Clustering

- Consider an enlarged optimization problem:

$$\min_{C, \{m_k\}_{k=1}^K} \sum_{k=1}^K \sum_{C(i)=k} \|x_i - m_k\|^2$$

- Hence, we are solving both
  - the cluster index function  $C(\cdot)$ ,
  - and also the cluster centers  $m_k, k = 1 \dots K$ .
- This problem is NP-hard for  $\geq 2$  dimensions.

# $K$ -means Clustering

- Combinatorial algorithm is too expansive.
- Instead, consider an algorithm that alternatively updates the two components:
  - $C$ , the cluster assignments
  - $\{m_k\}_{k=1}^K$ : the cluster means
- We will do **an iterative update** by:
  - 1) Fixing  $C$ , find the best  $\{m_k\}_{k=1}^K$
  - 2) Fixing  $\{m_k\}_{k=1}^K$ , find the best  $C$

# K-means Clustering

- **Fixing  $C$** , we know the cluster label of each subject. For any set  $\{i : C(i) = k\}$ , finding the mean is

$$m_k = \arg \min_m \sum_{C(i)=k} \|x_i - m\|^2.$$

- This is simply finding the mean within cluster  $k$ , i.e.

$$m_k = \frac{\sum_{C(i)=k} x_i}{\sum_i \mathbf{1}\{C(i) = k\}}$$

# *K*-means Clustering

- Fixing the cluster means  $\{m_k\}_{k=1}^K$ , to find the new cluster assignments, we simply recalculate the distance from an observation to each of the cluster mean.

$$C(i) = \arg \min_{k \in \{1, \dots, K\}} d(x_i, m_k)$$


- Hence each point will be assigned to the closest cluster mean




# $K$ -means Clustering

- A  $K$ -means Clustering algorithm:
  - 1) Randomly split the dataset into  $K$  different subsets. Assign each subsets a cluster label. Then iterate between 2) and 3).
  - 2) Given the cluster assignment  $C$ , calculate the cluster mean vectors  $m_1, \dots, m_K$ .
  - 3) Given the current set of means  $\{m_1, \dots, m_K\}$ , assign each observation to the closest current cluster mean.
- Stop the algorithm when  $C$  does not change


# Demonstration




# Demonstration




# Demonstration



# Demonstration



# Demonstration



# $K$ -means Clustering

- **Note:** We usually initiate the cluster labels randomly. However, this algorithm **does not guarantee to global minimum**. Example?
- The algorithm still has a **descent property**, which leads to a local minimizer.

# Alternative Version

- $K$ -medoids is an alternative version of  $K$ -means:
- Replace the second step by searching for the **one observation** that minimizes the distance to all others in the cluster

$$i_k^* = \arg \min_{i: C(i)=k} \sum_{C(i')=k} D(x_i, x_{i'})$$

- Use  $x_{i_k^*}$  as the “center” of cluster  $k$ .



# Applications

- See the supplementary [R](#) file
- Example 1: the [iris](#) data
- Example 2: cluster pixels in a picture

# Additional Issues

- Why do clustering when we have the true label?
- How to choose  $K$
- Which variables to include?

# Hierarchical Clustering


---

# Hierarchical Clustering

- Choosing the number of clusters  $K$  can be difficult
- A hierarchical representation which
  - at the lowest level, each cluster contains a single observation.
  - at the highest level there is only one cluster containing all observations.
- Use dendrogram to display the clustering result.


# Hierarchical Clustering

- Suppose we have a set of 6 observations




# Dendrogram


- A typical dendrogram from hierarchical clustering
- How to construct this?



# Demonstration




# Demonstration







# Demonstration



# Demonstration



# Demonstration



# Algorithm (agglomerative)

- An **agglomerative algorithm** is a “**bottom up**” approach:
  - Begin with every observation representing a singleton cluster.
  - At each step, merge two “closest” clusters into one cluster and reduce the number of clusters by one.
  - Stop when all observations are in the same cluster
- **How** to choose which two clusters to merge?
- This requires:
  - A distance measure between any **two observations**  $d(x_i, x_{j'})$
  - A distance measure between any **two sets of observations**  $d(G, H)$

# Distance Measures

- Distance  $d(G, H)$  between two clusters  $G$  and  $H$  can be defined in different ways:

- Complete linkage** (default of `hclust()`): the furthest pair

$$d(G, H) = \max_{i \in G, i' \in H} d_{ii'}$$

- Single linkage**: the closest pair

$$d(G, H) = \min_{i \in G, i' \in H} d_{ii'}$$

- Average linkage**: average dissimilarity

$$d(G, H) = \frac{1}{n_G n_H} \sum_{i \in G} \sum_{i' \in H} d_{ii'}$$

- Different choice may results in different hierarchical structures

# Distance Matrix

- To perform a hierarchical clustering, a matrix of all the pair-wise distances is sufficient
- We **don't have to know** the values of the original observations
- This is an  $n \times n$  matrix: the  $(i, i')$ 's element represents the distance between  $x_i$  and  $x_{i'}$
- This matrix is also called a **dissimilarity matrix**.
  - Symmetric
  - diagonal elements are zero

# Applications

- See the supplementary [R](#) file
- Example 1: the [iris](#) data
- Example 2: RNA expression data

# Principle Component Analysis

---



# Principle Component Analysis

- Principle Component Analysis (PCA) is an old but very useful technique invented by Karl Pearson in 1901
- The main purpose is data visualization (mostly in 2d)
- It also serves as a dimension reduction method
- Unsupervised method, can be used for preprocessing the data.

# Principle Component Analysis

- Given that we have a  $n \times p$  design matrix  $\mathbf{X}$ , there are many equivalent approaches (motivations):
  - **Explain the most variation:** Produce a derived set of uncorrelated variables  $\mathbf{Z}_k = \mathbf{X}\alpha_k$ ,  $k = 1, \dots, q < p$  that are linear combinations of the original variables, and explain most of the variation in the original set
  - **Approximate the design matrix:** Approximate the design matrix  $\mathbf{X}$  by the best (using Frobenius norm) rank- $q$  matrix, which can be performed through SVD

# Principle Component Analysis

- Suppose we have an  $n \times p$  design matrix  $\mathbf{X}$
- The first step is to **center each variable**, i.e., subtract the column means from each column respectively.
- In the following, we assume that  $\mathbf{X}$  is already centered.
- Centering  $\mathbf{X}$  does nothing but re-positioning the axis


# Singular Value Decomposition

- One way to understand the PCA is using a **singular value decomposition** (SVD)

$$\mathbf{X}_{n \times p} = \mathbf{U}_{n \times n} \mathbf{D}_{n \times p} \mathbf{V}_{p \times p}^T,$$

where both  $\mathbf{U}$  and  $\mathbf{V}$  are orthogonal matrices, i.e.

$\mathbf{U}^T \mathbf{U} = \mathbf{U} \mathbf{U}^T = \mathbf{I}$ , and same for  $\mathbf{V}$ ; and  $\mathbf{D}_{n \times p}$  is a diagonal matrix.



- The diagonal elements of  $\mathbf{D}_{n \times p}$  are of a decreasing order.

# Low Rank Approximation

- If we have to choose a rank-1 matrix  $\mathbf{A}$  to approximate  $\mathbf{X}_{n \times p}$ , what would we do?
- Turns out that the best choice is

$$\mathbf{U}_1 d_{11} \mathbf{V}_1^T,$$

where  $\mathbf{U}_1$  is the first column of  $\mathbf{U}$ ,  $\mathbf{V}_1$  is the first column of  $\mathbf{V}$ , and  $d_{11}$  is the first diagonal element of  $\mathbf{D}$

- In other words,  $\|\mathbf{X} - \mathbf{U}_1 d_{11} \mathbf{V}_1^T\|_2^2$  is minimized with this choice.
- Hence,  $\mathbf{U}_1 d_{11} \mathbf{V}_1^T$  is a rank-1 matrix that explained the variations in  $\mathbf{X}$  as much as possible.

# Principle Component Analysis

- Let's consider the sample covariance matrix  $\hat{\Sigma} = \mathbf{X}^T \mathbf{X} / (n - 1)$ , since  $\mathbf{X}$  is already centered.
- $\hat{\Sigma}$  can be diagonalize into

$$\hat{\Sigma} = \mathbf{V} \mathbf{D}^* \mathbf{V}^T,$$

where columns of  $\mathbf{V}$  are **principle directions** (loadings) and projecting  $\mathbf{X}$  on these loadings gives the **principal components**

- On the other hand, based on SVD,

$$\mathbf{X} = \mathbf{U} \mathbf{D} \mathbf{V}^T,$$

we can rewrite  $\hat{\Sigma}$  as

$$\hat{\Sigma} = \mathbf{V} \mathbf{D}^T \mathbf{U}^T \mathbf{U} \mathbf{D} \mathbf{V}^T / (n - 1) = \mathbf{V} \frac{\mathbf{D}^2}{n - 1} \mathbf{V}^T$$


# Principle Component Analysis

- So the right singular vectors  $\mathbf{V}$  of  $\mathbf{X}$  are just the principle directions, and the principal components are basically projecting each row (observation) of  $\mathbf{X}$  onto those directions:

$$\mathbf{XV} = \mathbf{UDV}^T\mathbf{V} = \mathbf{UD}$$


- The first column of  $\mathbf{U}$  is the first PC, and  $d_{11}$  is the variation along that direction, which is also the squared eigenvalue from SVD.
- PCA should be performed by **centering**  $\mathbf{X}$  first (column-wise, i.e., by each variable).

# Demonstration





# Demonstration



# Principle Component Analysis

- PCA is a **dimension reduction** tool, often used for visualization
- The leading PCs may display interesting information of the underlying (unobserved) clusters/manifold
- PCA is unsupervised, i.e., the directions does not necessarily reflect the relationship with the response (if there is any)