

STAT 432: Basics of Statistical Learning

Kernel Methods

Shiwei Lan, Ph.D. <shiwei@illinois.edu>

<http://shiwei.stat.illinois.edu/stat432.html>

March 28, 2019

University of Illinois at Urbana-Champaign

- K -Nearest Neighbor Revisited
- Kernel Smoother and Kernel Functions
- Local Polynomial Regression
- Kernel Density Estimation
- Multivariate Kernels

K-Nearest Neighbor Revisited

K-Nearest Neighbor

- *K*-nearest neighbor regression is weighted averaging
- Suppose we want to predict x . With training set $\{x_i, y_i\}_{i=1}^n$,

$$\hat{f}(x) = \sum_{i=1}^n w(x, x_i) y_i$$

where the weights


$$w(x, x_i) = \begin{cases} \frac{1}{k} & \text{if } x_i \in N_k(x) \\ 0 & \text{o.w.} \end{cases}$$

- $N_k(x)$ is a set of k observations in the neighborhood of x


K-Nearest Neighbor

- Problems:
 - Requires sorting after the weights are calculated — $\mathcal{O}(n \log(n))$
 - The weights $w(x, x_i)$ drop off abruptly to zero outside the neighborhood of x . This accounts for jagged appearance of the fit.
- To improve:
 - Easier method to assign weights
 - Smoothed weight functions

KNN vs. Gaussian Kernel



KNN vs. Gaussian Kernel



Kernel Smoother and Kernel Functions

Kernel Smoother (Univariate)

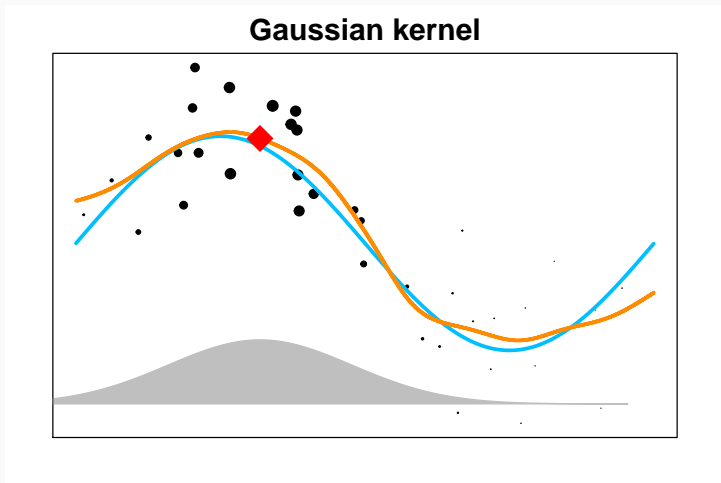
- We can still use the local averaging idea: Fit a simple model locally at each point x using only those observations close to it.
- Localization via the weighting function $K(x, x_i)$, the weight of x_i is based on its distance from x
- For any point x , we calculate the **weighted average**

$$\hat{f}(x) = \frac{\sum_i K_h(x, x_i) y_i}{\sum_i K_h(x, x_i)}$$

where h is a tuning parameter (called **bandwidth**) that controls the distance.

- The estimator is called **Nadaraya-Watson** kernel estimator
- $K_h(\cdot, \cdot)$ is a kernel function controlled by **bandwidth** h
 - Assigns larger value if the two inputs are closer to each other.
- Requires little or no training time. All the work gets done at evaluation time (same as k NN), however, no sorting is required.
- K NN is also a type of kernel method: the weight is $1/k$ or 0.

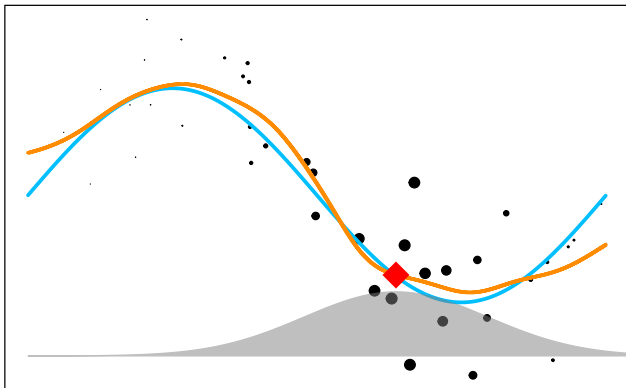
Gaussian Kernel



Predicting the target point $x = 2$

Gaussian Kernel

Gaussian kernel



Predicting the target point $x = 4$

Gaussian Kernel

- We often write the kernel function $K(x_1, x_2)$ in a different way, with $u = x_1 - x_2$

$$K_h(u) = h^{-1}K(u/h)$$

and $K(\cdot)$ is a “standard” version of the kernel (with $h = 1$).

- For example, a popular choice is the **Gaussian kernel**:

$$\begin{aligned} K(u) &= \phi(u) \\ &= \frac{1}{\sqrt{2\pi}} \exp \{ -u^2/2 \} \end{aligned}$$

- To incorporate the bandwidth (sd of Gaussian),

$$K_h(u) = \frac{1}{h\sqrt{2\pi}} \exp \{ -(u/h)^2/2 \}$$

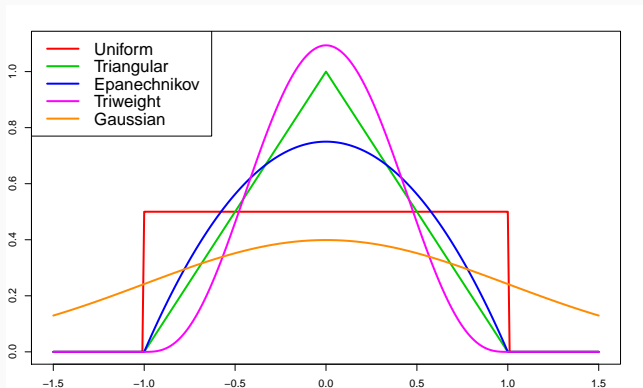
The Effect of Bandwidth

- The bandwidth h plays a crucial role:
- If h is large, we assign a relatively large kernel weight even when two points are far away from each other.
 - This works the same as a larger k in KNN
- If h is small, we assign a relatively small kernel weight as soon as two points are moving away from each other.
 - This works the same as a small k in KNN
- Large h : smoother estimate, bias \uparrow , variance \downarrow
- Small h : rougher estimate, bias \downarrow , variance \uparrow

Other Popular Kernels

- A kernel function usually satisfies the following properties:
 - K is properly normalized (e.g. pdf): $\int K(u)du = 1$;
 - K is symmetric around 0: $K(u) = K(-u)$;
 - $\int u^2 K(u)du \leq \infty$
 - $\int K^2(u)du \leq \infty$
- Many different kernel functions (besides Gaussian):
 - Uniform: $K(u) = \frac{1}{2} \cdot 1(|u| \leq 1)$
 - Triangular: $K(u) = (1 - |u|) \cdot 1(|u| \leq 1)$
 - Epanechnikov: $K(u) = \frac{3}{4}(1 - u^2) \cdot 1(|u| \leq 1)$
 - Triweight: $K(u) = \frac{35}{32}(1 - u^2)^3 \cdot 1(|u| \leq 1)$
 - ...
- They can all incorporate the bandwidth h

Different Kernel Functions



Different Kernel Functions

- So what's the difference between different kernels?
- **Efficiency** of a kernel function:
 - Efficiency is measured by $(\int u^2 K(u) du)^{\frac{1}{2}} \int K^2(u) du$
 - A quantity that evaluates the mean integrated squared error (MISE) of a kernel estimator defined as

$$\text{MISE}(\hat{f}) = \mathbb{E} \int (\hat{f}(x) - f(x))^2 dx$$


- Epanechnikov kernel is the most efficient.
- However, the **efficiency of other kernels** are not too bad: Gaussian is 95%; Uniform is 93% (the worst, relative to Epanechnikov kernel)
- **Choosing h is far more important than choosing the kernel.**

Local Linear Regression

Boundary Effects

- The Nadaraya-Watson kernel is notorious for boundary effects.
- In the previous Gaussian kernel example, there is a substantial bias at the boundaries.
- **Intuition**: all neighboring points are smaller/larger than the boundary $f(x)$.
- **Solution**: Locally weighted linear regression can make a first order correction (constants vs. straight lines)

Kernels



Local Linear Regression

- Recall that a (global) simple linear regression minimize the following objective function

$$\text{RSS} = \sum_{i=1}^n (y_i - \beta_0 - \beta_1 x_i)^2$$


- Idea: how about we pay more attention to the points that are closer to a target point x ?

RSS (locally weighted)

$$= \sum_{i=1}^n K_h(x, x_i) (y_i - \beta_0 - \beta_1 x_i)^2$$

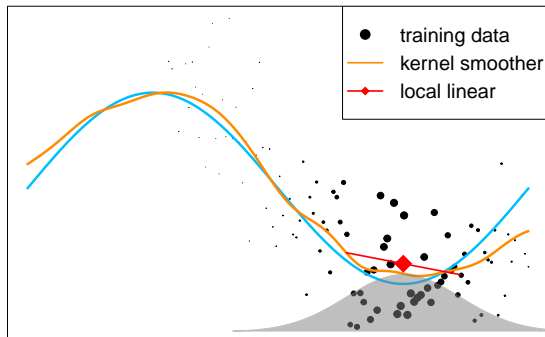
Local Linear Regression

Local Linear Regression at $x = 0$



Local Linear Regression

Local Linear Regression at $x = 4.712$



Local Linear Regression

- Similar to other kernel methods, the estimated (local) linear function is only valid at the local point x .
- The parameters β_0 and β_1 are in fact $\beta_0(x)$ and $\beta_1(x)$
- Hence, if we are interested in a **different target point**, we need to **refit the model entirely**. This could be computationally intense.
- The catch: more parameters to be estimated — smaller bias but larger variance

Local Polynomial Regression

- In general, we may consider a locally weighted d polynomial regression, which minimizes the local RSS objective function

$$\sum_{i=1}^n K_h(x, x_i) \left[y_i - \beta_0(x) - \sum_{r=1}^d \beta_r(x) x_i^r \right]^2$$

- Further reduces bias, at a price of higher variance.
- Warning: very sensitive to the choice of bandwidth h .
- Note: although its possible, but we **usually do not use** $r > 2$

Solving the Local Polynomial Regression

- Since the local polynomial regression is a weighted linear model, we may rewrite things in a matrix form:
- Let \mathbf{W} be a $n \times n$ diagonal matrix defined as

$$\mathbf{W} = \text{diag}(K_h(x, x_1), K_h(x, x_2), \dots, K_h(x, x_n))$$

- Then the weighted RSS can be written as

$$\sum_{i=1}^n K_h(x, x_i) (y_i - \beta_0 - \beta_1 x_i)^2 = (\mathbf{y} - \mathbf{X}\boldsymbol{\beta})^\top \mathbf{W} (\mathbf{y} - \mathbf{X}\boldsymbol{\beta})$$

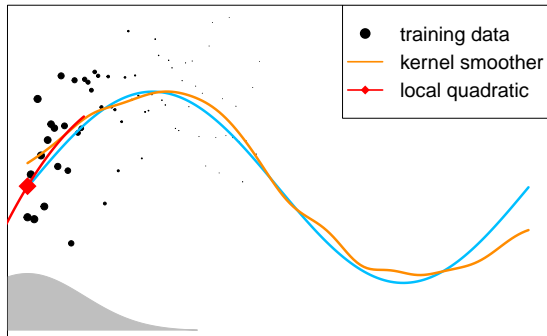
- And the solution is (from normal equations)

$$\hat{\boldsymbol{\beta}} = (\mathbf{X}^\top \mathbf{W} \mathbf{X})^{-1} \mathbf{X}^\top \mathbf{W} \mathbf{y}$$

- Note that we need to recalculate this for each target points x .


Local Quadratic Regression

Local Quadratic Regression at $x = 0$



Local Quadratic Regression

Local Quadratic Regression at $x = 4.712$



R implementations

- R function `loess` provides fitting of the local polynomial regressions
- The most important parameter `span` = α controls the degree of smoothing: only αn number of closest points are used based on the distance $|x - x_i|$, forming the neighborhood “ $N(x)$ ”
- A weighted least-square linear regression is fit within the neighborhood
- The weights uses tri-cube kernel: $w_{x,i} = (1 - u^3)^3$ with

$$u_i = \frac{|x_i - x|}{\max_{N(x)} |x_j - x|}$$

- `degree` specifies the degree of the polynomial
- Other implementations such as `locfit` and `locpoly` (use Gaussian kernel)

Kernel Density Estimation

Kernel Density Estimation

- Another area where we often use the kernel methods is **estimating the density**
- Given some observations from an unknown distribution, we want to estimate the pdf of that distribution (unsupervised)

$$X_1, \dots, X_n \stackrel{\text{i.i.d}}{\sim} f(\cdot)$$

- Some density estimation methods
 - Histograms
 - Assume a family of distributions and estimate parameters
 - Kernel density estimator

Histogram Estimator of Density Functions

- If $f(\cdot)$ is the pdf of X , then we have:

$$\int f(u)du = 1, \quad \text{and} \quad f(u) > 0 \quad \text{for all } u$$

$$\int_{x-\frac{h}{2}}^{x+\frac{h}{2}} f(u)du = \mathbf{P} \left(x - \frac{h}{2} \leq X \leq x + \frac{h}{2} \right)$$


$$f(x) = \lim_{h \rightarrow 0} \frac{1}{h} \mathbf{P} \left(x - \frac{h}{2} \leq X \leq x + \frac{h}{2} \right)$$

- A natural estimator, with a set of observations $\{x_i\}_{i=1}^n$, is

$$\hat{f}(x) = \frac{\#\{x_i : x_i \in [x - \frac{h}{2}, x + \frac{h}{2}]\}}{hn}$$

- This is very similar to the **histogram estimator**. But it is bumpy and non-smooth.

Uniform Kernel Density Estimation



Kernel Density Estimation

- Parzen estimate

$$\hat{f}(x) = \frac{1}{n} \sum_{i=1}^n K_h(x, x_i)$$

- $K_h(\cdot)$ here is a kernel function, controlled by h .
- Again, we can use the popular Gaussian kernel function for this task

$$K_h(x, x_i) = \frac{1}{h\sqrt{2\pi}} \exp \left\{ -\frac{(x - x_i)^2}{2h^2} \right\}$$

- `hist` makes histograms
- `density` for kernel density estimator
- `bw.nrd` and a set of related functions for bandwidth selection
- The rule of thumb (Silverman 1986) for h in univariate case is

$$\hat{h} = 1.06\hat{\sigma}n^{-1/5}$$

where $\hat{\sigma}$ is the sample standard deviation.

Multivariate Kernel Estimations

Multivariate Density Functions

- We can extend the idea from univariate to multivariate case.
- The only thing that needs to be changed is the kernel function:

$$K_{\mathbf{H}}(\mathbf{u}, \mathbf{v})$$

for any p dimensional vectors \mathbf{u} and \mathbf{v} , and a kernel bandwidth matrix \mathbf{H} .

- If we are still using a Gaussian density function, then

$$K_{\mathbf{H}}(\mathbf{u}, \mathbf{v}) = \frac{1}{(2\pi)^{p/2} |\mathbf{H}|^{1/2}} \exp \left\{ -\frac{1}{2} (\mathbf{u} - \mathbf{v})^{\top} \mathbf{H}^{-1} (\mathbf{u} - \mathbf{v}) \right\}$$

- A simplified version is to take \mathbf{H} as a diagonal matrix.